

YOLO Reproduction-4

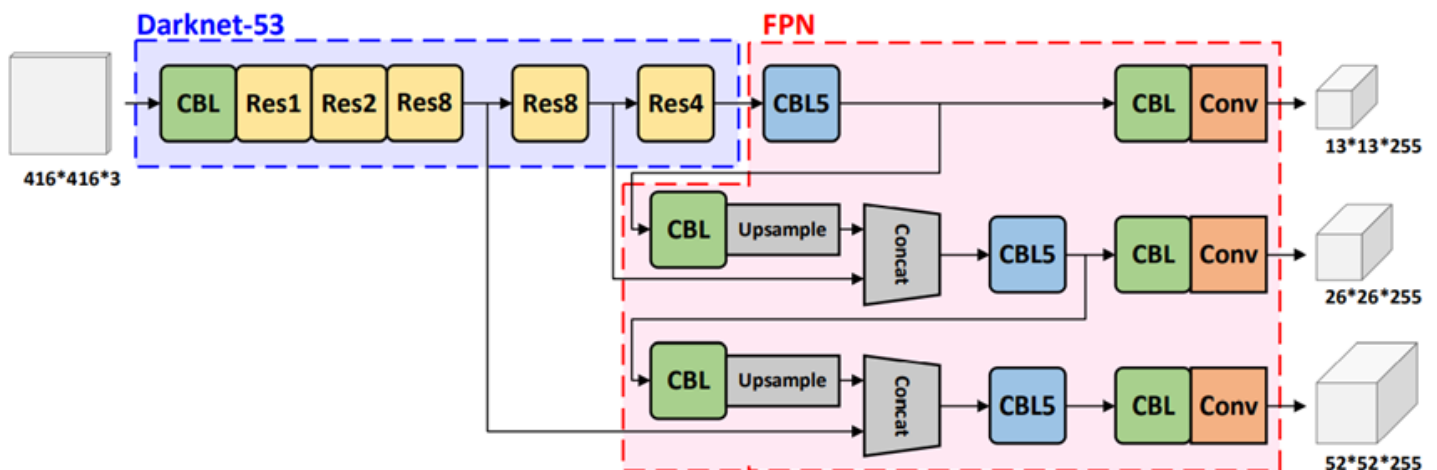
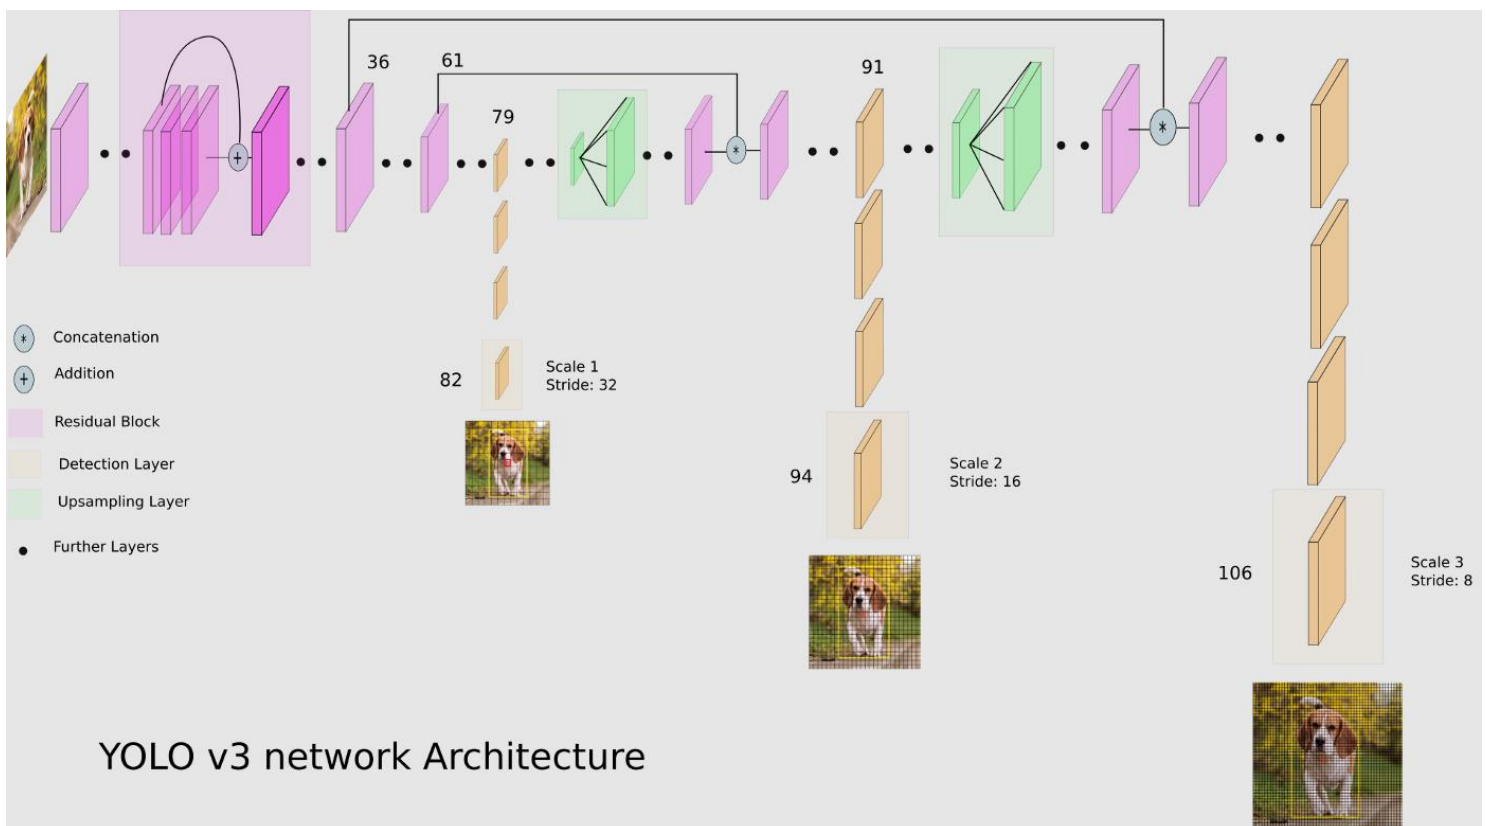
Advisor: Dr. Chih-Yu Wang

Presenter: Shao-Heng Chen

Date: July, 27, 2022

1. YOLOv3 problems

- (1) currently still untrainable
- (2) original YOLOv3 network architecture



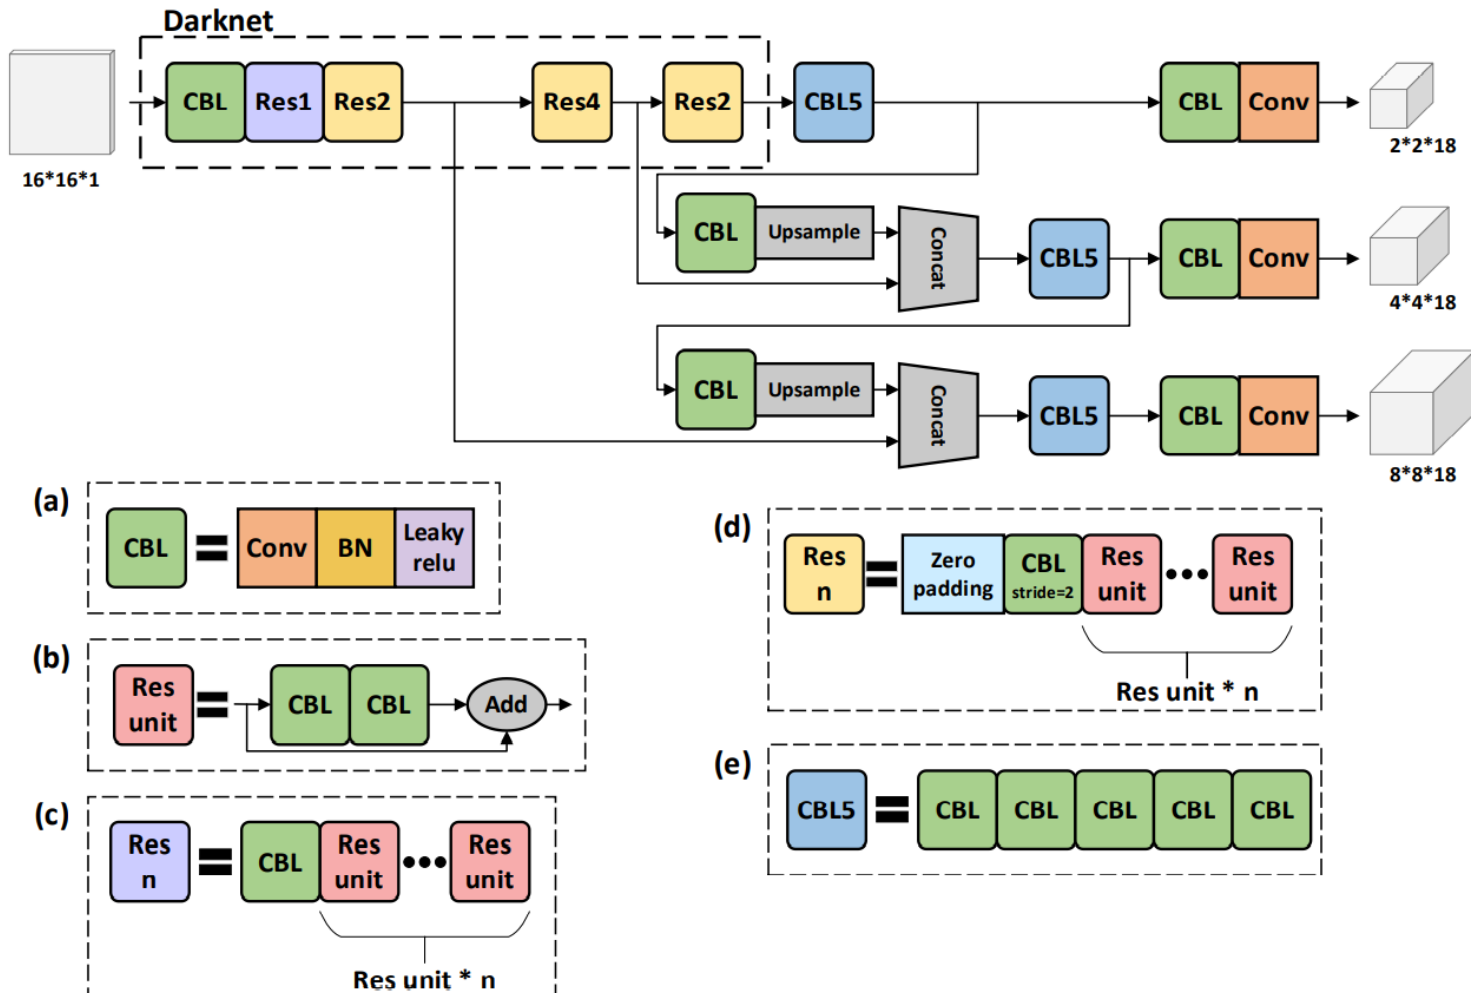
(3) YOLO-CFAR network architecture

- YOLO-CFAR vs. Keras YOLOv3 model comparison

(https://github.com/paulchen2713/YOLO_project/commit/ae46523c274b97774db01dd9af90bc8c48dc174f)

- YOLOv3-PyTorch model

(https://github.com/paulchen2713/YOLO_project/commit/05fe39a7036da9ff71c32b6f027ab93d8490379b)



```
# -*- coding: utf-8 -*-
"""
Created on Mon Jul 18 17:04:43 2022
```

```
@author: Paul
```

```
@file: model.py
```

```
@dependencies:
```

```
env pt3.7
```

```
python 3.7.13
```

```
torch >= 1.7.1
```

```
torchvision >= 0.8.2
```

```
@references:
```

```
Redmon, Joseph and Farhadi, Ali, YOLOv3: An Incremental Improvement, April 8, 2018.
```

```
(https://doi.org/10.48550/arXiv.1804.02767)
```

Ayoosh Kathuria, Whats new in YOLO v3?, April, 23, 2018. (<https://towardsdatascience.com/yolo-v3-object-detection-53fb7d3bfe6b>)

Sanna Persson, YOLOv3 from Scratch, Mar 21, 2021. (<https://sannaperzon.medium.com/yolov3-implementation-with-training-setup-from-scratch-30ecb9751cb0>)

Implementation of YOLOv3 architecture

"""

```
import torch
```

```
import torch.nn as nn
```

"""

Information about architecture config:

 Tuple is structured by (filters, kernel_size, stride)

 Every conv is a same convolution.

 List is structured by "B" indicating a residual block followed by the number of repeats

 "S" is for scale prediction block and computing the yolo loss

 "U" is for upsampling the feature map and concatenating with a previous layer

"""

```
config = [
```

```
    (32, 3, 1), # (32, 3, 1) is the CBL, CBL = Conv + BN + LeakyReLU
```

```
    (64, 3, 2),
```

```
    ["B", 1], # (64, 3, 2) + ["B", 1] is the Res1, Res1 = ZeroPadding + CBL + (CBL + CBL + Add)*1
```

```
    (128, 3, 2),
```

```
    ["B", 2], # (128, 3, 2) + ["B", 2] is th Res2, Res2 = ZeroPadding + CBL + (CBL + CBL + Add)*2
```

```
    (256, 3, 2),
```

```
    ["B", 8], # (256, 3, 2) + ["B", 8] is th Res8, Res8 = ZeroPadding + CBL + (CBL + CBL + Add)*8
```

```
    (512, 3, 2),
```

```
    ["B", 8], # (512, 3, 2) + ["B", 8] is th Res8, Res8 = ZeroPadding + CBL + (CBL + CBL + Add)*8
```

```
    (1024, 3, 2),
```

```
    ["B", 4], # (1024, 3, 2) + ["B", 4] is th Res4, Res4 = ZeroPadding + CBL + (CBL + CBL + Add)*4
```

to this point is Darknet-53 which has 52 layers

52 = 1 + (1 + 1*2) + (1 + 2*2) + (1 + 8*2) + (1 + 8*2) + (1 + 4*2) ?

```
    (512, 1, 1), #
```

```
    (1024, 3, 1), #
```

```
    "S",
```

```
    (256, 1, 1),
```

```
    "U",
```

```
    (256, 1, 1),
```

```
    (512, 3, 1),
```

```
    "S",
```

```
    (128, 1, 1),
```

```
    "U",
```

```
    (128, 1, 1),
```

```
    (256, 3, 1),
```

```

"S",
# 252 = 1 + 3 + (4+7) + (4+7*2) + (4+7*8) + (4+7*8) + (4+7*4) + 19 + 5 + 19 + 5 + 19 ?
]

```

```

config = [
    (32 // 2, 3, 1),
    (64 // 2, 3, 2),
    ["B", 1],      # (64, 3, 2) + ["B", 1] is the Res1
    (128, 3, 2),
    ["B", 2],      # (128, 3, 2) + ["B", 2] is th Res2
    # (256, 3, 2),
    # ["B", 8],      # (256, 3, 2) + ["B", 8] is th Res8
    (512, 3, 2),
    ["B", 4],      # (512, 3, 2) + ["B", 8] is th Res8
    (1024 // 2, 3, 2),
    ["B", 1],      # ["B", 4], to this point is Darknet-53, which has 53 layers?
    # 52 = 1 + (1 + 1*2) + (1 + 2*2) + (1 + 8*2) + (1 + 8*2) + (1 + 4*2) ?
    (512 // 2, 1, 1),
    (1024, 3, 1),
    "S",
    (256, 1, 1),
    "U",
    (256 // 2, 1, 1),
    (512 // 2, 3, 1),
    "S",
    (128 // 2, 1, 1), #
    "U",
    (128 // 2, 1, 1),
    (256 // 2, 3, 1),
    "S",
    # 252 = 1 + 3 + (4+7) + (4+7*2) + (4+7*8) + (4+7*8) + (4+7*4) + 19 + 5 + 19 + 5 + 19 ?
]

```

```

class CNNBlock(nn.Module):
    def __init__(self, in_channels, out_channels, bn_act=True, **kwargs):
        super(CNNBlock, self).__init__()
        # if we do use bn activation function in the block, then we do not want to use bias, its unnecessary
        # **kwargs will be the kernal size, the stride and padding as well
        self.conv = nn.Conv2d(in_channels, out_channels, bias=not bn_act, **kwargs)
        self.bn = nn.BatchNorm2d(out_channels)
        self.leaky = nn.LeakyReLU(negative_slope=0.1) # default negative_slope=0.01

```

```
self.use_bn_act = bn_act # indicating if the block is going to use a batch norm NN activation function
```

```
def forward(self, x):
```

```
    # using if-else statement in the forward pass might lose on some performance, negligible?
```

```
    # we use bn activation by default, except for scale prediction
```

```
    if self.use_bn_act:
```

```
        return self.leaky(self.bn(self.conv(x))) # bn_act()
```

```
    # for scale prediction we don't want to use batch norm LeakyReLU on our output, just normal Conv
```

```
    else:
```

```
        return self.conv(x)
```

```
class ResidualBlock(nn.Module):
```

```
    def __init__(self, channels, use_residual=True, num_repeats=1):
```

```
        super(ResidualBlock, self).__init__()
```

```
        self.layers = nn.ModuleList()
```

```
        for _ in range(num_repeats): # repeat for num_repeats
```

```
            self.layers += [
```

```
                nn.Sequential(
```

```
                    CNNBlock(channels, channels // 2, kernel_size=1, padding=0), # down samples or reduces the number
```

of filters

```
                    # CNNBlock(channels // 2, channels, kernel_size=3, padding=1), # then brings it back again
```

```
                    CNNBlock(channels // 2, channels, kernel_size=3, padding=1),
```

```
                )
```

```
            ]
```

```
    # 1. why specify use_residual in a ResidualBlock? is because in some cases we are going to use skip
```

```
    # connections, in some cases we just going through the config file and build the ordinary ResidualBlock
```

```
    # 2. why we need to store these?
```

```
    self.use_residual = use_residual # indicating using residual
```

```
    self.num_repeats = num_repeats # number of repeats set to 1 by default
```

```
def forward(self, x):
```

```
    for layer in self.layers:
```

```
        x = layer(x) + x if self.use_residual else layer(x)
```

```
    # if self.use_residual:
```

```
        # x = x + layer(x)
```

```
        # x = layer(x) + x
```

```
    # else:
```

```
        # x = layer(x)
```

```
    return x
```

```

class ScalePrediction(nn.Module):
    def __init__(self, in_channels, num_classes):
        super(ScalePrediction, self).__init__()

        # for every single cell grid we have 3 anchor boxes, for every anchor box we have 1 node for each of the
        # classes

        # for each bounding box we have [P(Object), x, y, w, h] and that's 5 values
        self.pred = nn.Sequential(
            # CNNBlock(in_channels, 2 * in_channels, kernel_size=3, padding=1),
            CNNBlock(in_channels, 2 * in_channels, kernel_size=3, padding=1),
            CNNBlock(2 * in_channels, 3 * (num_classes + 5), bn_act=False, kernel_size=1),
        )
        self.num_classes = num_classes

    def forward(self, x):
        # we want to return the prediction of x, then we want to reshape it to the number of examples in our batch
        # split out_channel "3 * (num_classes + 5)" into two different dimensions "3, (num_classes + 5)", instead of
        # having a long vector of bounding boxes, and change the order of the dimensions
        return (
            self.pred(x)
            .reshape(x.shape[0], 3, self.num_classes + 5, x.shape[2], x.shape[3])
            .permute(0, 1, 3, 4, 2)
        )
        # [x.shape[0], 3, x.shape[2], x.shape[3], self.num_classes + 5], e.g. [N, 3, 13, 13, 5+num_classes]
        # for scale one, we have N examples in our batch, each example has 3 anchors, each anchor has 13-by-13 grid
        # and every cell has (5+num_classes) output, output dimension = N x 3 x 13 x 13 x (5+num_classes)

class YOLOv3(nn.Module):
    def __init__(self, in_channels=3, num_classes=1):
        super(YOLOv3, self).__init__()
        self.num_classes = num_classes
        self.in_channels = in_channels

        # we want to create the layers using the config file, and store them in a nn.ModuleList()
        self.layers = self._create_conv_layers() # we immediately call _create_conv_layers() to initialize the layers

    def forward(self, x):
        # need to keep track of outputs and route connections
        outputs = [] # we have one output for each scale prediction, should be 3 in total
        route_connections = [] # e.g. after upsampling, we concatenate the channels of skip connections

        for i, layer in enumerate(self.layers):
            if isinstance(layer, ScalePrediction):

```

```

        outputs.append(layer(x)) # we're going to add that layer
        continue # and then continue from where we were previously, not after ScalePrediction

    # calling layer(x) is equivalent to calling layers.__call__(x), and __call__() is actually calling
    layer.forward(x)

    # which is defined in class layer(nn.Module), but in practice we should use layer(x) rather than
    layer.forward(x)

    x = layer(x) #
    print(f"layer {i}: ", x.shape)

    # skip layers are connected to ["B", 8] based on the paper, original config file
    if isinstance(layer, ResidualBlock) and layer.num_repeats != 1: #
    # if isinstance(layer, ResidualBlock) and layer.num_repeats == 8:
        route_connections.append(x)

    elif isinstance(layer, nn.Upsample): # if we use the Upsample
        # we want to concatenate with the last route connection, with the last one we added
        x = torch.cat([x, route_connections[-1]], dim=1) # why concatenate along dimension 1 for the channels
        route_connections.pop() # after concatenation, we remove the last one

    # print(f"outputs: {outputs}")
    return outputs

# create the layers using the config files
def _create_conv_layers(self):
    layers = nn.ModuleList() # keep track of all the layers in a ModuleList, which supports tools like
model.eval()

    in_channels = self.in_channels # only need to specify the first in_channels, I suppose

    # go through and parse the config file and construct the model line by line
    for module in config:
        # if it's a tuple (filters, kernel_size, stride), e.g. (32, 3, 1), then it's just a CNNBlock
        if isinstance(module, tuple):
            out_channels, kernel_size, stride = module # we want to take out the (filters, kernel_size, stride)
            layers.append(
                CNNBlock(
                    in_channels,
                    out_channels,
                    kernel_size=kernel_size,
                    stride=stride,
                    # padding=1 if kernel_size == 3 else 0, # if kernel_size == 1 then padding = 0
                    padding=1 if kernel_size == 3 else 0,

```

```

    )
)
# the in_channels for the next block is going to be the out_channels of this block
in_channels = out_channels # update the in_channels of the next layer

# if it's a List, e.g. ["B", 1], then it's a ResidualBlock
elif isinstance(module, List):
    num_repeats = module[1] # we want to take out the number of repeats, which is going to be module[1]
    # and module[0] should be "B", which indicates that this is a ResidualBlock
    layers.append(ResidualBlock(in_channels, num_repeats=num_repeats,))

# if it's a String, e.g. "S" or "U", then it might be ScalePrediction or Upsampling
elif isinstance(module, str):
    # "S" for ScalePrediction
    if module == "S":
        layers += [
            ResidualBlock(in_channels, use_residual=False, num_repeats=1),
            CNNBlock(in_channels, in_channels // 2, kernel_size=1),
            ScalePrediction(in_channels // 2, num_classes=self.num_classes),
        ]
        # after ScalePrediction, we want to continue from CNNBlock, since we have scale_factor=2
        in_channels = in_channels // 2 # we then want to divide in_channels by 2
    # "U" for Upsampling
    elif module == "U":
        layers.append(nn.Upsample(scale_factor=2,))
        in_channels = in_channels * 3 # 3 == 2 + 1, concatenated the channels from previously

return layers

```

```

if __name__ == "__main__":
    # actual parameters
    num_classes = 1 # 20
    # YOLOv1: 448, YOLOv2/YOLOv3: 416 (with multi-scale training)
    IMAGE_SIZE = 16 # multiples of 32 are workable with stride [32, 16, 8]
    # stride = [8, 4, 2]
    stride = [16, 8, 4] # 16
    # stride = [32, 16, 8] # 32

    # simple test settings
    num_examples = 2
    num_channels = 3 # num_anchors

```



```

model = YOLOv3(num_classes=num_classes) # initialize a YOLOv3 model as model
# simple test with random inputs of 2 examples, 3 channels, and IMAGE_SIZE-by-IMAGE_SIZE input
x = torch.randn((num_examples, num_channels, IMAGE_SIZE, IMAGE_SIZE))
out = model(x)

print("Output Shape: ")
print("[num_examples, num_channels, feature_map, feature_map, num_classes + 5]")
for i in range(num_channels):
    print(out[i].shape)

assert out[0].shape == (2, 3, IMAGE_SIZE//stride[0], IMAGE_SIZE//stride[0], num_classes + 5) # [2, 3, 13, 13,
num_classes + 5]
assert out[1].shape == (2, 3, IMAGE_SIZE//stride[1], IMAGE_SIZE//stride[1], num_classes + 5) # [2, 3, 26, 26,
num_classes + 5]
assert out[2].shape == (2, 3, IMAGE_SIZE//stride[2], IMAGE_SIZE//stride[2], num_classes + 5) # [2, 3, 52, 52,
num_classes + 5]
print("Success!")

# layer 0: torch.Size([2, 16, 16, 16])
# layer 1: torch.Size([2, 32, 8, 8])
# layer 2: torch.Size([2, 32, 8, 8])
# layer 3: torch.Size([2, 128, 4, 4])
# layer 4: torch.Size([2, 128, 4, 4])
# layer 5: torch.Size([2, 512, 2, 2])
# layer 6: torch.Size([2, 512, 2, 2])
# layer 7: torch.Size([2, 512, 1, 1])
# layer 8: torch.Size([2, 512, 1, 1])
# layer 9: torch.Size([2, 256, 1, 1])
# layer 10: torch.Size([2, 1024, 1, 1])
# layer 11: torch.Size([2, 1024, 1, 1])
# layer 12: torch.Size([2, 512, 1, 1])
# layer 14: torch.Size([2, 256, 1, 1])
# layer 15: torch.Size([2, 256, 2, 2])
# layer 16: torch.Size([2, 128, 2, 2])
# layer 17: torch.Size([2, 256, 2, 2])
# layer 18: torch.Size([2, 256, 2, 2])
# layer 19: torch.Size([2, 128, 2, 2])
# layer 21: torch.Size([2, 64, 2, 2])
# layer 22: torch.Size([2, 64, 4, 4])
# layer 23: torch.Size([2, 64, 4, 4])
# layer 24: torch.Size([2, 128, 4, 4])

```

```
# layer 25: torch.Size([2, 128, 4, 4])
# layer 26: torch.Size([2, 64, 4, 4])
# Output Shape:
# [num_examples, num_channels, feature_map, feature_map, num_classes + 5]
# torch.Size([2, 3, 1, 1, 6])
# torch.Size([2, 3, 2, 2, 6])
# torch.Size([2, 3, 4, 4, 6])
# Success!

# layer 0: torch.Size([2, 32, 416, 416])
# layer 1: torch.Size([2, 64, 208, 208])
# layer 2: torch.Size([2, 64, 208, 208])
# layer 3: torch.Size([2, 128, 104, 104])
# layer 4: torch.Size([2, 128, 104, 104])
# layer 5: torch.Size([2, 256, 52, 52])
# layer 6: torch.Size([2, 256, 52, 52])
# layer 7: torch.Size([2, 512, 26, 26])
# layer 8: torch.Size([2, 512, 26, 26])
# layer 9: torch.Size([2, 1024, 13, 13])
# layer 10: torch.Size([2, 1024, 13, 13])
# layer 11: torch.Size([2, 512, 13, 13])
# layer 12: torch.Size([2, 1024, 13, 13])
# layer 13: torch.Size([2, 1024, 13, 13])
# layer 14: torch.Size([2, 512, 13, 13])
# layer 16: torch.Size([2, 256, 13, 13])
# layer 17: torch.Size([2, 256, 26, 26])
# layer 18: torch.Size([2, 256, 26, 26])
# layer 19: torch.Size([2, 512, 26, 26])
# layer 20: torch.Size([2, 512, 26, 26])
# layer 21: torch.Size([2, 256, 26, 26])
# layer 23: torch.Size([2, 128, 26, 26])
# layer 24: torch.Size([2, 128, 52, 52])
# layer 25: torch.Size([2, 128, 52, 52])
# layer 26: torch.Size([2, 256, 52, 52])
# layer 27: torch.Size([2, 256, 52, 52])
# layer 28: torch.Size([2, 128, 52, 52])
# Output Shape:
# [num_examples, num_channels, feature_map, feature_map, num_classes + 5]
# torch.Size([2, 3, 13, 13, 6])
# torch.Size([2, 3, 26, 26, 6])
# torch.Size([2, 3, 52, 52, 6])
# Success!
```